

# HarvardX: PH125.9x Data Science: Capstone Course

## Movie Rating Prediction Project; Movie Lens Dataset

*Bijal Ashwin Ved 12/28/2019*

### Contents

|  |           |
|--|-----------|
| <b>1 Introduction</b>                      | <b>2</b>  |
| 1.1 Dataset . . . . .                      | 2         |
| <b>2 Analysis and Results</b>              | <b>9</b>  |
| 2.1 Simple Prediction Based on Mean Rating | 9         |
| 2.2 Movie Effects Model . . . . .          | 9         |
| 2.3 Movie and User Effects Model . . . . . | 10        |
| 2.4 Regularization. . . . .                | 11        |
| <b>3 Results and Discussion</b>            | <b>13</b> |
| <b>4 Conclusion</b>                        | <b>13</b> |

# 1 Introduction

The Movie Lens project forms part of the HarvardX: PH125.9x Data Science: Capstone course. It closely follows the Data Science textbook by Rafael Irizarry particularly; Chapter 34.7 Recommendation systems, and Chapter 34.9 Regularization. The dataset was supplied as part of the course.

The aim of the project is to develop and train a recommendation machine learning algorithm to predict a rating given by a users to a movie in the dataset. The Residual Mean Square Error (RMSE) will be used to evaluate the accuracy of the algorithm. The required criteria for the projects is a RMSE lower than 0.8775.

This report will present an overview of the data, analysis, results and a conclusion.

## 1.1 Dataset

An introductory review of the dataset is performed in order to familiarise ourselves. Data is downloaded as per instruction from the MovieLens 10M dataset.

```
#####  
## Create edx set, validation set  
#####  
  
# Note: this process could take a couple of minutes  
  
# Package Installs  
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")  
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")  
  
# Libraries  
library(tidyverse)  
library(caret)  
library(hexbin)  
  
# MovieLens 10M dataset:  
# https://grouplens.org/datasets/movielens/10m/  
# http://files.grouplens.org/datasets/movielens/ml-10m.zip  
  
# Check for download  
datafile <- "MovieLens.RData"  
if(!file.exists("MovieLens.RData"))  
{  
  print("Download")  
  dl <- tempfile()  
  download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)  
  
  ratings <- read.table(text = gsub(":", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),  
                        col.names = c("userId", "movieId", "rating", "timestamp"))  
  
  movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\:", 3)  
  colnames(movies) <- c("movieId", "title", "genres")  
  movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],  
                                             title = as.character(title),  
                                             genres = as.character(genres))  
  
  movielens <- left_join(ratings, movies, by = "movieId")  
  
# Validation set will be 10% of MovieLens data
```

```

set.seed(1, sample.kind = "Rounding") # if using R 3.6.0: set.seed(1, sample.kind = "Rounding")
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set

validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set

removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)
rm(dl, ratings, movies, test_index, temp, movielens, removed)
save(edx, validation, file = datafile)

} else {
  load(datafile)
}

```

Edx dataset contains rows corresponding to an users rating of a movie. The set contains the variables; "userId", "movieId", "rating", "timestamp", "title", "genres".

```

# Summarise Data
head(edx, 5)

```

```

##   userId movieId rating timestamp                title
## 1      1     122      5 838985046      Boomerang (1992)
## 2      1     185      5 838983525      Net, The (1995)
## 4      1     292      5 838983421      Outbreak (1995)
## 5      1     316      5 838983392      Stargate (1994)
## 6      1     329      5 838983392 Star Trek: Generations (1994)
##                                     genres
## 1                      Comedy|Romance
## 2          Action|Crime|Thriller
## 4 Action|Drama|Sci-Fi|Thriller
## 5          Action|Adventure|Sci-Fi
## 6 Action|Adventure|Drama|Sci-Fi

```

Summarising the dataset reveals a well formatted set with no missing values. Movies are rated between 0.5 and 5.0, with 9000055 rows in total.

```
summary(edx)
```

```

##      userId      movieId      rating      timestamp
## Min.   : 1  Min.   : 1  Min.   :0.500  Min.   :7.897e+08
## 1st Qu.:18124 1st Qu.: 648 1st Qu.:3.000 1st Qu.:9.468e+08
## Median :35738 Median : 1834 Median :4.000 Median :1.035e+09
## Mean   :35870 Mean   : 4122 Mean   :3.512 Mean   :1.033e+09
## 3rd Qu.:53607 3rd Qu.: 3626 3rd Qu.:4.000 3rd Qu.:1.127e+09
## Max.   :71567 Max.   :65133 Max.   :5.000 Max.   :1.231e+09
##      title      genres
## Length:9000055 Length:9000055

```

```
## Class :character    Class :character
## Mode  :character    Mode  :character
##
##
##
```

The dataset contains ~10,700 unique movies, ~70,000 unique movies, and ~800 unique combinations of genres, and a mean movie rating of ~3.5 out of 5.

```
# Movies, Users and Genres in Database
edx %>% summarise(
  uniq_movies = n_distinct(movieId),
  uniq_users = n_distinct(userId),
  uniq_genres = n_distinct(genres))
```

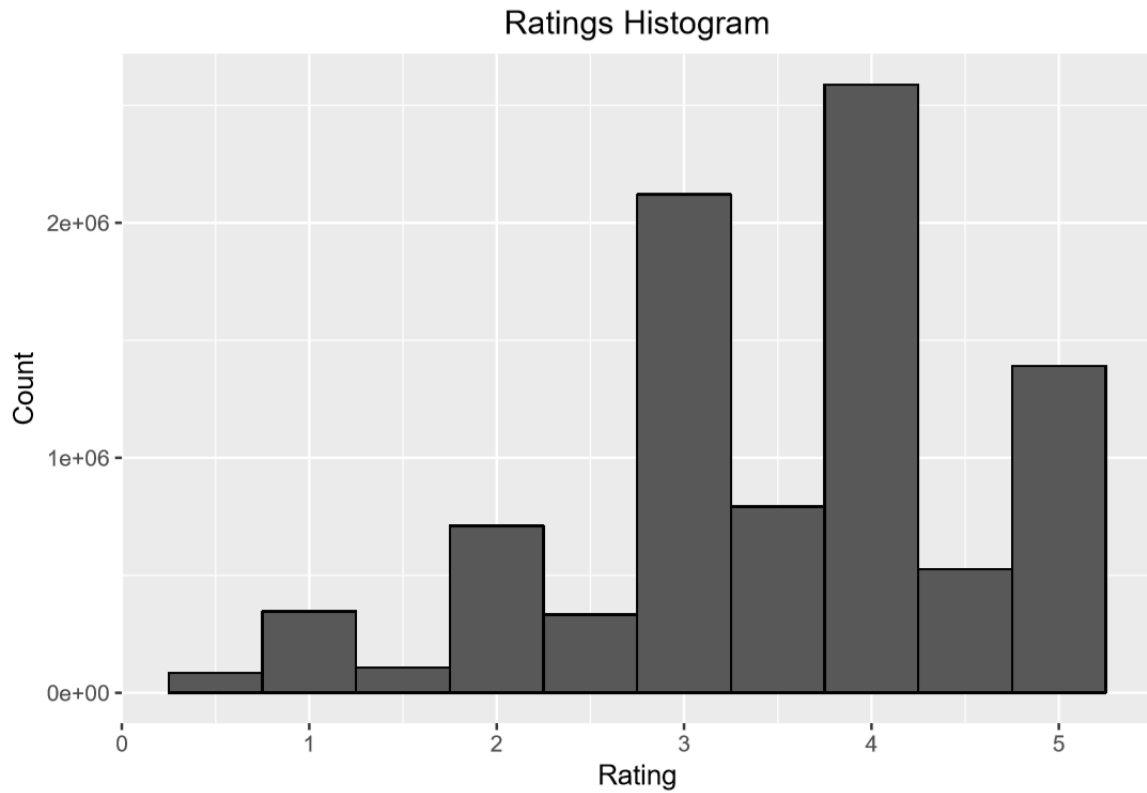
```
##   uniq_movies uniq_users uniq_genres
## 1         10677      69878         797
```

```
# Ratings Mean
rating_mean <- mean(edx$rating)
rating_mean
```

```
## [1] 3.512465
```


A histogram of the dataset mapping ratings and counts are below. The highest rating users give movies are, in decreasing order 4.0, 3.0, 5.0. Overall users rate movies higher in the ratings scale. Similarly, users are more likely to rate full ratings rather than half ratings on the scale.

```
# Ratings Histogram
edx %>%
  ggplot(aes(rating)) +
  geom_histogram(binwidth = 0.5, color = "black") +
  xlab("Rating") +
  ylab("Count") +
  ggtitle("Ratings Histogram") +
  theme(plot.title = element_text(hjust = 0.5))
```



The majority of users rate between 10 and 100 movies, whilst some may rate over 1,000. Including a variable in the model to account for number of ratings should be discussed.


```
# Ratings Users - Number of Ratings
edx %>%
  count(userId) %>%
  ggplot(aes(n)) +
  geom_histogram(color = "black", bins=30) +
  scale_x_log10() +
  xlab("# Ratings") +
  ylab("# Users") +
  ggtitle("Users Number of Ratings") +
  theme(plot.title = element_text(hjust = 0.5))
```



A heatmap showing average movie rating against number of movies rated is plotted. The most common movies ratings and number of movies rated are highlighted. This occurs for a rating between 3.5 and 4.0 and between 10 and 100 movies.


A linear curve was fitted to the data to show the overall trajectory of the ratings with number of movies rated. A more memory and calculation intensive curve fitting was attempted but unfortunately unable to run on the author's computer. The more ratings a user gives results in a lower mean rating.

```
# Ratings Users - Mean by Number with Curve Fitted
edx %>%
  group_by(userId) %>%
  summarise(mu_user = mean(rating), number = n()) %>%
  ggplot(aes(x = mu_user, y = number)) +
  geom_bin2d( ) +
  scale_fill_gradientn(colors = grey.colors(10)) +
  labs(fill="User Count") +
  scale_y_log10() +
  geom_smooth(method = lm) +
  ggtitle("Users Average Ratings per Number of Rated Movies") +
  xlab("Average Rating") +
  ylab("# Movies Rated") +
  theme(plot.title = element_text(hjust = 0.5))
```



The number of ratings for each movie are shown below in the histogram. A number of outlier movies have been rated less than 10 times which will make predicting future ratings more difficult.

```
# Ratings Movies - Number of Ratings
edx %>%
  count(movieId) %>%
  ggplot(aes(n)) +
  geom_histogram(color = "black", bins=30) +
  scale_x_log10() +
  xlab("# Ratings") +
  ylab("# Movies") +
  ggtitle("Movie Number of Ratings") +
  theme(plot.title = element_text(hjust = 0.5))
```





## 2 Analysis and Results

The Residual Mean Square Error (RMSE) is the error function to that will measure accuracy and quantify the typical error we make when predicting the movie rating. An error larger than 0.8775, it means our typical error is larger than the required for this assignment almost a star, which is not good. RMSE defined;

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

where; N is the number of users, movie ratings, and the sum incorporating the total combinations.

### 2.1 Simple Prediction Based on Mean Rating

The simple prediction model uses the mean of the dataset to predict the rating for all movies. The model assumes that all differences are due to a random error;

$$Y_{u,i} = \mu + \epsilon_{u,i}$$

where  $Y_{u,i}$  is the prediction,  $\epsilon_{u,i}$  is the independent error, and  $\mu$  the expected “true” rating for all movies. Predicting the mean gives the following naive RMSE.

```
## Simple Prediction based on Mean Rating
mu <- mean(edx$rating)
mu

## [1] 3.512465

rmse_naive <- RMSE(validation$rating, mu)
rmse_naive

## [1] 1.061202

## Save Results in Data Frame
rmse_results = data_frame(method = "Naive Analysis by Mean", RMSE = rmse_naive)

## Warning: `data_frame()` is deprecated, use `tibble()`.
## This warning is displayed once per session.

rmse_results %>% knitr::kable()
```

| method                 | RMSE     |
|------------------------|----------|
| Naive Analysis by Mean | 1.061202 |

Investigating the dataset allows for more advanced analysis and rating predictions with smaller error.

### 2.2 Movie Effects Model

The Movie Effects Model calculates a bias term for each movie based on the difference between the movies mean rating and the overall mean rating.

$$Y_{u,i} = \mu + b_i + \epsilon_{u,i}$$

where  $Y_{u,i}$  is the prediction,  $\epsilon_{u,i}$  is the independent error, and  $\mu$  the mean rating for all movies, and  $b_i$  is the bias for each movie  $i$ .

```
## Simple model taking into account the movie effects, b_i

mu <- mean(edx$rating)
mu

## [1] 3.512465

movie_avgs <- edx %>%
  group_by(movieId) %>%
  summarise(b_i = mean(rating - mu))

predicted_ratings <- mu + validation %>%
  left_join(movie_avgs, by='movieId') %>%
  pull(b_i)

rmse_model_movie_effects <- RMSE(predicted_ratings, validation$rating)
rmse_model_movie_effects

## [1] 0.9439087

rmse_results <- bind_rows(rmse_results,
  data_frame(method="Movie Effects Model",
    RMSE = rmse_model_movie_effects))
rmse_results %>% knitr::kable()
```

| method                 | RMSE      |
|------------------------|-----------|
| Naive Analysis by Mean | 1.0612018 |
| Movie Effects Model    | 0.9439087 |

The Movie Effect Model; predicting the movie rating with both bias,  $b_i$ , and mean,  $\mu$  gives an improved prediction with a lower RMSE value.

### 2.3 Movie and User Effects Model

The next step is to incorporate the individual User Effects,  $b_u$ , in to the model. Acknowledging each user inherent bias to mark all films higher or lower.

$$Y_{u,i} = \mu + b_i + b_u + \epsilon_{u,i}$$

where  $Y_{u,i}$  is the prediction,  $\epsilon_{u,i}$  is the independent error, and  $\mu$  the mean rating for all movies,  $b_i$  is the bias for each movie  $i$ , and  $b_u$  is the bias for each user  $u$ .

```
## Movie and User Effects Model
# Simple model taking into account the user effects, b_u

user_avgs <- edx %>%
  left_join(movie_avgs, by="movieId") %>%
  group_by(userId) %>%
  summarise(b_u = mean(rating - mu - b_i))

predicted_ratings <- validation %>%
  left_join(movie_avgs, by='movieId') %>%
  left_join(user_avgs, by='userId') %>%
  mutate(pred = mu + b_i + b_u) %>%
  pull(pred)
```

```
rmse_model_user_effects <- RMSE(predicted_ratings, validation$rating)
rmse_model_user_effects

## [1] 0.8653488

rmse_results <- bind_rows(rmse_results,
                          data_frame(method="Movie and User Effects Model",
                                     RMSE = rmse_model_user_effects))
rmse_results %>% knitr::kable()
```

| method                       | RMSE      |
|------------------------------|-----------|
| Naive Analysis by Mean       | 1.0612018 |
| Movie Effects Model          | 0.9439087 |
| Movie and User Effects Model | 0.8653488 |

Incorporating the user bias into the model resulted in a further reduced RMSE.

## 2.4 Regularisation

Regularisation allows for reduced errors caused by movies with few ratings which can influence the prediction and skew the error metric. The method uses a tuning parameter,  $\lambda$ , to minimise the RMSE. Modifying  $b_i$  and  $b_u$  for movies with limited ratings.

$$Y_{u,i} = \mu + b_{i,n,\lambda} + b_{u,n,\lambda} + \epsilon_{u,i}$$

```
# Predict via regularisation, movie and user effect model
# (as per https://rafalab.github.io/dsbook/34.9 Regularization)

lambdas <- seq(0, 10, 0.25)
rmse <- sapply(lambdas, function(l){

  mu <- mean(edx$rating)

  b_i <- edx %>%
    group_by(movieId) %>%
    summarise(b_i = sum(rating - mu)/(n()+1))

  b_u <- edx %>%
    left_join(b_i, by="movieId") %>%
    group_by(userId) %>%
    summarise(b_u = sum(rating - b_i - mu)/(n()+1))

  predicted_ratings <- validation %>%
    left_join(b_i, by = "movieId") %>%
    left_join(b_u, by = "userId") %>%
    mutate(pred = mu + b_i + b_u) %>%
    pull(pred)

  return(RMSE(predicted_ratings, validation$rating))


})
```

```
rmse_regularisation <- min(rmses)
rmse_regularisation
```

```
## [1] 0.864817
```

```
# Plot RMSE against Lambdas to find optimal lambda
```

```
qplot(lambdas, rmses)
```



```
lambda <- lambdas[which.min(rmses)]
lambda
```

```
## [1] 5.25
```

```
rmse_results <- bind_rows(rmse_results,
  data_frame(method="Regularised Movie and User Effects Model",
    RMSE = rmse_regularisation))
```

```
rmse_results %>% knitr::kable()
```

| method                                   | RMSE      |
|--|-----------|
| Naive Analysis by Mean                   | 1.0612018 |
| Movie Effects Model                      | 0.9439087 |
| Movie and User Effects Model             | 0.8653488 |
| Regularised Movie and User Effects Model | 0.8648170 |

Regulisation of a Movie and User Effect model has lead to a lowest RMSE of the prediction methods for the MovieLens ratings system.

### 3 Results and Discussion

The final values of the prediction models are shown below;

```
rmse_results %>% knitr::kable()
```

| method                                   | RMSE      |
|--|-----------|
| Naive Analysis by Mean                   | 1.0612018 |
| Movie Effects Model                      | 0.9439087 |
| Movie and User Effects Model             | 0.8653488 |
| Regularised Movie and User Effects Model | 0.8648170 |

The models from most accurate to least accurate are as follows; Regularised Movie and User Effects Model; Movie and User Effects Model; Movie Effects Model; and Simple Average Model.

The final model optimised for the prediction is the following;

$$Y_{u,i} = \mu + b_{i,n,\lambda} + b_{u,n,\lambda} + \epsilon_{u,i}$$

The lowest value of RMSE predicted is 0.8648170.

### 4 Conclusion

A machine learning algorithm to predict the ratings from the Movie Lens dataset was constructed. The optimal model incorporated the effects of both user and movie bias were incorporated in the model, and these variables were regularised to incorporate movies with a low number of ratings.

The aim of the project was to develop an algorithm lower than 0.87750, which was achieved by the Movie and User Effects and Regularised Movie and User Effects model.